

Foundations of High Performance Computing

Lecture 7: Domain Decomposition by means of MPI



2021-2022 Stefano Cozzini

“Foundation of HPC” course
DATA SCIENCE &
SCIENTIFIC COMPUTING

Agenda

- Domain Decomposition approach
- A Real example

Domain Decomposition

- A fundamental technique in both parallel programming and MPI programming
- The way to solve very large problems in memory on distributed memory architecture
- Useful to know some trick of the trade to deal with that using MPI

What is domain decomposition?

- Perform a numerical simulation (e.g. solve an equation) on a domain of given shape
- Discretize the domain using a mesh or particles
- Split the work (and memory!) among N processors
- Can be done “geographically” or purely on the data, irrespective of their spatial location
- Domains can have complex shapes

Goals to achieve..


- All processors having equal (in terms of wall- clock time) amounts of work
==> “Load balancing”
- Minimum number of communication steps between the processors
==> “Communication scheduling”
- Minimum amount of data that needs to be communicated. No deadlocks.

Domain Decomposition

Cartesian

0	1	2	6	7	8
3	4	5	9	10	11
12	13	14	18	19	20
15	16	17	21	22	23

Unstructured



Examples with 4 sub-domains

Non geographical approaches


- Partition the data irrespective of their physical location in the computational domain
- Example: long range MD
 - having N interacting particles, each processor receives N/P particles. Since all of them have to communicate to all anyway, neighborhood relations are not important
 - Best load balance due to finest granularity
 - Easy to compute
 - No re-balancing needed if particles move

Cartesian coordinates
=
regular meshes

Operations on meshes


The values at mesh nodes are updated such as to numerically solve a given equation:

- **Finite differences** (Stencil methods)
Values of mesh nodes are updated using the values of (any set of) nearby nodes.
- Ex: a 5-point stencil (diffusion equation)




Operations on meshes

- Ghost layers around each sub-domain provide the needed values at the sub-domain boundaries




Ghost cells or halo area

- To calculate the new grid points of a sub-domain, additional grid points from other sub-domains are needed.
- They are stored in halos (ghost cells, shadows)
- Halo depends on form of stencil




1. Splitting the domain..

- Consider a 3D box
- We can split along 1/2/3 dimensions
- Which is the best choice in term of performance (i.e scalability) ?



Example: a cube


- Let us consider the following parameters:
 - N^3 size of the cubic matrix
 - P processor
 - Cyclic/Periodic boundary condition:
 - **two neighbors** in each direction
 - W : width of the “ghost cells” or halo
- Define in term of these parameters the cost of the communication message



Domain decomposition: 1D

- Splitting along 1 dimension:


1D communication
 $= n^2 * \mathbf{2} * w * \mathbf{1}$



Domain decomposition: 2D

- Splitting along 2 dimension:


$$\begin{aligned} &\mathbf{2D} \text{ communication} \\ &= n^2 * \mathbf{2} * w * \mathbf{2/P^{1/2}} \end{aligned}$$



Domain decomposition: 3D

- Splitting along 3 dimension:

$$\text{3D communication} = n^2 * \textcolor{red}{2} * w * \textcolor{blue}{3/P^{2/3}}$$



What about scalability of communication ?

Process	1D	2D	3D
			0
2,00	2,00	2,83	3,78
4,00	2,00	2,00	2,38
8,00	2,00	1,41	1,50
16,00	2,00	1,00	0,94
32,00	2,00	0,71	0,60
64,00	2,00	0,50	0,38

Domain Decomposition 101

1. Split the domain into blocks.
2. Assign blocks to MPI-processes one-to-one.
3. Provide a "map" of neighbors to each process.
4. Write or modify your code so it only updates a single block.
5. Insert communication subroutine calls where needed.
6. Adjust the boundary conditions code.
7. Use "ghost cells/halo".


Real example: Jacobi solver..

- prototype for many stencil-based iterative methods in numerical analysis and simulation
- Basic form: solve the diffusion equation for a scalar function: $\Phi(\mathbf{r}, t)$

$$\frac{\partial \Phi}{\partial t} = \Delta \Phi,$$

Straightforward 2D serial implementation: the stencil

```
do k = 1,kmax
  do i = 1,imax
    phi(i,k,t1) = 0.25 * phi(i+1,k,t0) + 0.25 * phi(i-1,k,t0)
                  + 0.25 * phi(i,k+1,t0) + 0.25 * phi(i,k-1,t0)
  enddo
enddo
```



All taken from reference 4

The full serial algorithm..


- Ensure that the code produces a converged result.

```
double precision, dimension(0:imax+1,0:kmax+1,0:1) :: phi
double precision :: maxdelta,eps
integer :: t0,t1
eps = 1.d-14 ! convergence threshold
t0 = 0 ; t1 = 1
maxdelta = 2.d0*eps
do while(maxdelta.gt.eps)
  maxdelta = 0.d0
  do k = 1,kmax
    do i = 1,imax
      phi(i,k,t1) = 0.25 * phi(i+1,k,t0) + 0.25 * phi(i-1,k,t0)
                  + 0.25 * phi(i,k+1,t0) + 0.25 * phi(i,k-1,t0)
      maxdelta = max(maxdelta,abs(phi(i,k,t1)-phi(i,k,t0)))
    enddo
  enddo
  ! swap arrays
  i = t0 ; t0=t1 ; t1=i
enddo
```

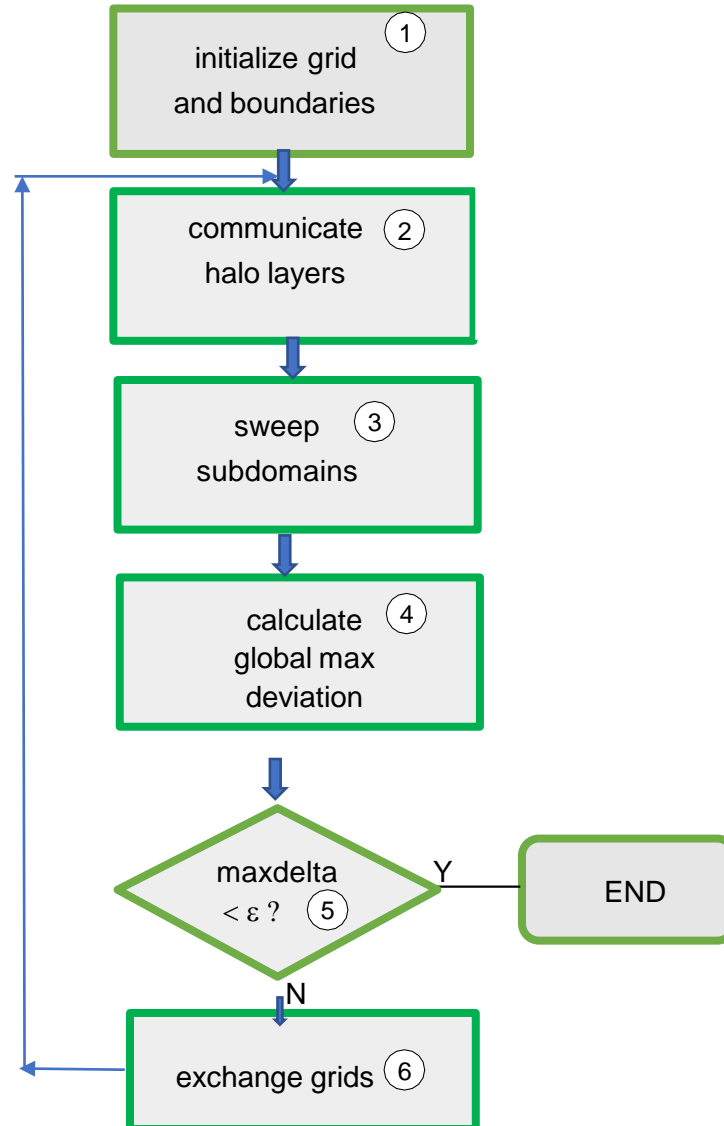
What happens in the parallel implementation ?

- The computational core works the same but:
 - Convergence criterion is no longer enough on subdomains but need to be computed globally
 - ➔ requires a reduce operation among all processors
 - We need to take care of boundary conditions : Cell close to border require special care and require halo layers

The parallel data distribution..



The parallel algorithm



A few remarks

- Initialization is done with virtual topology


```
1 call MPI_Dims_create(numprocs, 3, proc_dim, ierr)
2
3 if(myid.eq.0) write(*, '(a,3(i3,x))') 'Grid: ', &
4     (proc_dim(i), i=1,3)
5
6 l_reorder = .true.
7 call MPI_Cart_create(MPI_COMM_WORLD, 3, proc_dim, pbc_check, &
8     l_reorder, GRID_COMM_WORLD, ierr)
9
10 if(GRID_COMM_WORLD .eq. MPI_COMM_NULL) goto 999
11
12 call MPI_Comm_rank(GRID_COMM_WORLD, myid_grid, ierr)
13 call MPI_Comm_size(GRID_COMM_WORLD, nump_grid, ierr)
```

Halo communication

- We use point to point communication to exchange halo layer.
- Point-to-point communication requires consecutive message buffers.
- Do we have contiguous location in memory for the halo ?

Not at all...

- only those boundary cells that are consecutive in the inner (i) dimension are also consecutive in memory (fortran column-major order..)
- Whole layers in the i-j, i-k, and j-k planes are never consecutive

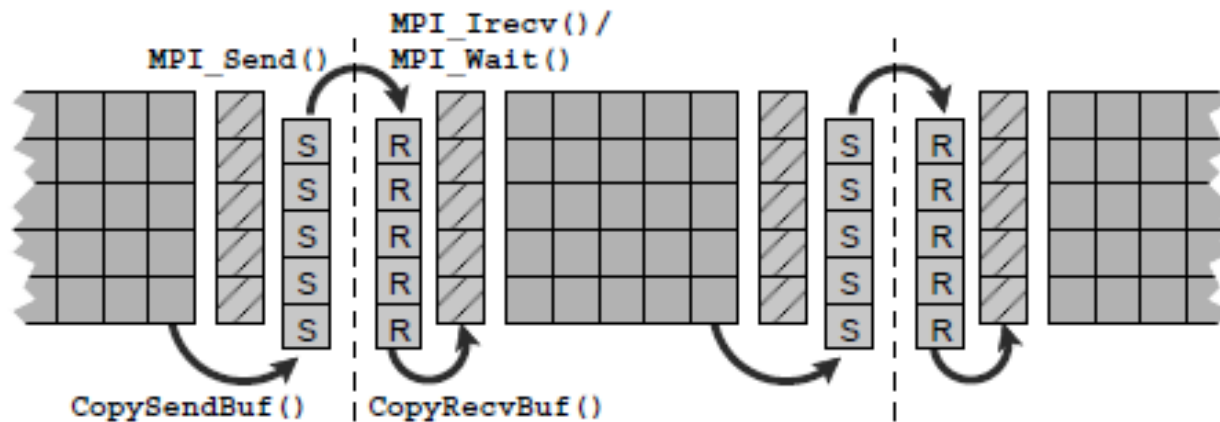


How to deal with this issue ?

- Send each chunk in a separate message
discouraged:
 - This approach would flood the network with short messages, and latency has to be paid for every request
- **BEST SOLUTION**
 - intermediate buffer must be used to gather boundary data to be communicated to a neighbor's ghost layer.

Halo exchange in one direction..

- We use two intermediate buffers per process, one for sending and one for receiving.



Dimension of the buffer...

- halo data can be different along different Cartesian directions
 - ➔ the size of the intermediate buffer must be chosen to accommodate the largest possible halo

```
1 integer, dimension(1:3) :: totmsgsize
2
3 ! j-k plane
4 totmsgsize(3) = loca_dim(1)*loca_dim(2)
5 MaxBufLen=max(MaxBufLen,totmsgsize(3))
6 ! i-k plane
7 totmsgsize(2) = loca_dim(1)*loca_dim(3)
8 MaxBufLen=max(MaxBufLen,totmsgsize(2))
9 ! i-j plane
10 totmsgsize(1) = loca_dim(2)*loca_dim(3)
11 MaxBufLen=max(MaxBufLen,totmsgsize(1))
12
13 allocate(fieldSend(1:MaxBufLen))
14 allocate(fieldRecv(1:MaxBufLen))
```

Halo exchange.

```

4  do disp = -1, 1, 2
5    do dir = 1, 3
6
7      call MPI_Cart_shift(GRID_COMM_WORLD, (dir-1), &
8                          disp, source, dest, ierr)
9
10     if(source /= MPI_PROC_NULL) then
11       call MPI_Irecv(fieldRecv(1), totmsgsize(dir), &
12                     MPI_DOUBLE_PRECISION, source, &
13                     tag, GRID_COMM_WORLD, req(1), ierr)
14     endif    ! source exists
15
16     if(dest /= MPI_PROC_NULL) then
17       call CopySendBuf(phi(iStart, jStart, kStart, to), &
18                       iStart, iEnd, jStart, jEnd, kStart, kEnd, &
19                       disp, dir, fieldSend, MaxBufLen)
20
21       call MPI_Send(fieldSend(1), totmsgsize(dir), &
22                     MPI_DOUBLE_PRECISION, dest, tag, &
23                     GRID_COMM_WORLD, ierr)
24     endif    ! destination exists
25
26     if(source /= MPI_PROC_NULL) then
27       call MPI_Wait(req, status, ierr)
28
29       call CopyRecvBuf(phi(iStart, jStart, kStart, to), &
30                       iStart, iEnd, jStart, jEnd, kStart, kEnd, &
31                       disp, dir, fieldRecv, MaxBufLen)
32     endif    ! source exists
33
34   enddo    ! dir
35 enddo    ! disp
36
```

For full 3D implementation

- See github repo...

Exercise

- Familiarize with the program and start running it on different number of processors and different sizes.
- Tomorrow we discuss performance model and scalability behavior